



Budapest University of Technology and Economics
Department of Measurement and Information Systems
Fault Tolerant Systems Research Group

Service Integration course

Cassandra

Oszkár Semeráth

Gábor Szárnyas

August 11, 2014

Contents

1	Installation	2
1.1	Linux	2
1.2	Windows	2
1.3	Documentation pointers	3
2	Laboratory exercises	4
2.1	Database design	4
2.2	Java client	7
2.3	Integration	11
2.4	Version control	13
2.5	Eclipse tips	14

Chapter 1

Installation

Apache Cassandra (<http://cassandra.apache.org/>) is an open source NoSQL database management system based on the column family data model.



Figure 1.1: The logo of Apache Cassandra

1.1 Linux

```
sudo mkdir /var/lib/cassandra
sudo mkdir /var/log/cassandra
sudo chown -R $USER:$USER /var/lib/cassandra
sudo chown -R $USER:$USER /var/log/cassandra
```

1.2 Windows

Set the JAVA_HOME environment variable to to the JDK's installation directory (Cassandra will work with the JRE as well, but some other Java-based applications – such as Maven – will not).

Install **Cygwin** with **Python 2.7.x** (if you want to add the package to an existing Cygwin installation, you have to re-run the installer). Navigate to `/cygdrive/c/apache-cassandra-2.0.5`. Run the shell with the following command:

```
$ bin/cqlsh
Connected to Test Cluster at localhost:9160.
[cqlsh 4.1.1 | Cassandra 2.0.5 | CQL spec 3.1.1 | Thrift protocol 19.39.0]
Use HELP for help.
```

Note that Cassandra itself is not supported under Cygwin (<http://stackoverflow.com/questions/12355507/error-when-starting-cassandra-with-bin-cassandra-f>).

1.3 Documentation pointers

- Official Apache wiki (<http://wiki.apache.org/cassandra/>). It is not very informative, use the DataStax documentation (<http://www.datastax.com/docs>) instead.
- CQL documentation: <http://www.datastax.com/documentation/cql/3.1/index.html>

Chapter 2

Laboratory exercises

2.1 Database design

1. You may always reset the database with the following command:

```
DROP KEYSPACE bicycle_database;
```

2. Create the **keyspace**:

```
CREATE KEYSPACE bicycle_database  
WITH replication = {  
  'class': 'SimpleStrategy',  
  'replication_factor': 1  
};
```

3. Create a **column family**:

```
USE bicycle_database;  
  
CREATE TABLE bicycles (  
  manufacturer text,  
  model text,  
  PRIMARY KEY (manufacturer, model)  
);  
  
INSERT INTO bicycles (manufacturer, model) VALUES ('KTM', 'Strada 1000');
```

4. Run a simple query:

```
SELECT * FROM bicycles;
```

5. Add some columns.

```
ALTER TABLE bicycles ADD year int;  
ALTER TABLE bicycles ADD size int;
```

6. Insert a few lines.

```
INSERT INTO bicycles (manufacturer, model, year, size) VALUES ('KTM', 'Strada 1000', 2014, 57);  
INSERT INTO bicycles (manufacturer, model, year, size) VALUES ('Trek', '1.2', 2013, 54);
```

7. Run a simple query again:

```
cqlsh:bicycle_database> SELECT * FROM bicycles;
```

manufacturer	model	size	year
KTM	Strada 1000	57	2014
Trek	1.2	54	2013

(2 rows)

8. We would like to collect the bicycles with the frame size of 54:

```
SELECT * FROM bicycles WHERE size = 54;
```

This throws the following error: Bad Request: No indexed columns present in by-columns clause with Equal operator.

9. Create indices.

```
CREATE INDEX bicycles_size ON bicycles (size);  
CREATE INDEX bicycles_year ON bicycles (year);
```

10. Re-run the query. It may not return any results as the rows were added *before* the index was created. Add the rows again and the query will work fine.

11. Try a query which filters for inequality:

```
SELECT * FROM bicycles WHERE year > 2013;
```

12. This also throws an error:

Bad Request: No indexed columns present in by-columns clause with Equal operator

13. Add an equality filter as well:

```
SELECT * FROM bicycles WHERE year > 2013 AND size = 57;
```

This time, the error message is the following: Bad Request: Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING

14. Allow filtering:

```
SELECT * FROM bicycles WHERE year > 2013 AND size = 57 ALLOW FILTERING;
```

This query works just fine. For details, see <http://wiki.apache.org/cassandra/SecondaryIndexes>.

15. Try using an OR logical operation in the query:

```
SELECT * FROM bicycles WHERE size = 54 OR size = 57;
```

The parser does not even recognize the operation: Bad Request: line 1:39 missing EOF at 'OR'.

16. CQL3 supports collections – let's use them for storing the number of cogs:

```
ALTER TABLE bicycles ADD front_cogs list<int>;  
ALTER TABLE bicycles ADD rear_cogs list<int>;
```




Figure 2.1: Drivetrain. Source: http://commons.wikimedia.org/wiki/File:Derailleur_Bicycle_Drivetrain.svg

- Update the KTM bicycle with an INSERT operation (source: http://www.ktm-bikes.at/bikes/road.html?action=bike_details&bike_id=160&cHash=dbfc043b8dd2263b610187cf9d093fc7).

```
INSERT INTO bicycles (manufacturer, model, front_cogs, rear_cogs)
VALUES ('KTM', 'Strada 1000', [50, 34], [12, 13, 14, 15, 17, 19, 21, 23, 25, 28]);
```

Query it, the result should look like this:

manufacturer	model	front_cogs	rear_cogs	size	year
KTM	Strada 1000	[50, 34]	[12, 13, 14, 15, 17, 19, 21, 23, 25, 28]	57	2014
Trek	1.2	null	null	54	2013

2.2 Java client

- Create a new **Maven project**. Choose the `org.apache.maven.archetypes:maven-archetype-quickstart` archetype. Set the **Group Id** to `hu.bme.mit.sysint` and the **Artifact Id** to `client`.
- Go to the `App` class and run it.
- Go to the `pom.xml` file. The `Ctrl+F11` hotkey does not run the application if we are on the POM file. Solution: go to **Window | Preferences | Run/Debug | Launching**, in the **Launch Operation** groupbox choose **Always launch the previously launched application**.
- Add the DataStax driver dependency (<https://github.com/datastax/java-driver>):

```
<dependency>
  <groupId>com.datastax.cassandra</groupId>
  <artifactId>cassandra-driver-core</artifactId>
  <version>2.0.0</version>
</dependency>
```

- Create the `CassandraClient` class.

```
private static final String KEYSPACE_NAME = "bicycle_database";
protected Session session;
protected Cluster cluster;

public CassandraClient(String host) {
    cluster = Cluster.builder().addContactPoints(host).build();
    session = cluster.connect("system");
}
```

- Instantiate the `CassandraClient` class from the main method:

```
public class App {
    private static final String CASSANDRA_HOST = "127.0.0.1";

    public static void main(String[] args) {
        client = new CassandraClient(CASSANDRA_HOST);
        client.load();
    }
}
```

- The following warning is shown.


```
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
```

8. Pick a logger implementation, e.g. **log4j** and search for its Maven dependency:

```
<dependency>
  <groupId>org.slf4j</groupId>
  <artifactId>slf4j-log4j12</artifactId>
  <version>1.7.6</version>
</dependency>
```

9. Run the project.

10. The following warning is displayed:

```
log4j:WARN No appenders could be found for logger (com.datastax.driver.core.FrameCompressor).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
```

11. Create log4j.properties file in the src/main/java folder. The content should look something like this:

```
# Root logger option
log4j.rootLogger=INFO, stdout

# Direct log messages to stdout
log4j.appender.stdout=org.apache.log4j.ConsoleAppender
log4j.appender.stdout.Target=System.out
log4j.appender.stdout.layout=org.apache.log4j.PatternLayout
log4j.appender.stdout.layout.ConversionPattern=%d{yyyy-MM-dd HH:mm:ss} %-5p %c{1}:%L - %m%n
```

Source: <http://www.mkyong.com/logging/log4j-log4j-properties-examples/>

12. If you'd like to hide the warning messages (e.g. for a production system), change the value of the log4j.rootLogger property from INFO to ERROR.
13. Alternatively, we may add the missing libraries:

```
<dependency>
  <groupId>org.xerial.snappy</groupId>
  <artifactId>snappy-java</artifactId>
  <version>1.1.0.1</version>
</dependency>

<dependency>
  <groupId>net.jpountz.lz4</groupId>
  <artifactId>lz4</artifactId>
  <version>1.2.0</version>
</dependency>
```

14. Use the following commands in the CQL shell:

```
DESCRIBE KEYSPACE bicycle_database;
DESCRIBE TABLE bicycles;
```

15. The application does not stop after it finishes its operations. This is because it has open connections. To solve this, implement the `Closeable` interface in the `CassandraClient` class:

```
public class CassandraClient implements Closeable {

    [...]

    public void close() {
        session.close();
        cluster.close();
    }
}
```

Java 7 has a feature similar to the `using` block in C#.

```
public static void main(String[] args) {
    try (CassandraClient client = new CassandraClient(CASSANDRA_HOST)) {
        client.load();
    } finally {
    }
}
```

16. We start the load with dropping the keyspace.

```
session.execute("DROP KEYSPACE bicycle_database");
```

This works fine for the first time, however, it throws an exception if the keyspace does not exist:

Exception in thread "main" com.datastax.driver.core.exceptions.InvalidQueryException: Cannot drop non existing keyspace 'bicycle_database'.

17. While the Javadoc of the `Session` interface's `execute()` method shows that it may throw an exception, the application is not required to catch this exception as it is a `RuntimeException`. For details, see the Javadoc of the `RuntimeException` class.

```
try {
    session.execute("DROP KEYSPACE bicycle_database");
} catch (InvalidQueryException e) {
    System.out.println("Cannot drop keyspace.");
}
```

A more elegant solution is to use:

```
DROP KEYSPACE IF EXISTS bicycle_database;
```

18. Create the keyspace:

```
String createKeyspace = "CREATE KEYSPACE " + KEYSPEC_NAME
    + " WITH replication = {\r\n"
    + "   'class': 'SimpleStrategy',\r\n"
    + "   'replication_factor': '1'\r\n" + "};";
session.execute(createKeyspace);
session.execute("USE " + KEYSPEC_NAME);
```

19. Create the table. Autoformatting may be messy here. To correct this, Go to **Window | Preferences | Java | Code Style | Formatter**. Create a new profile (e.g. Eclipse FTSG) and edit it. Go to the **Off/On Tags** tab and check **Enable Off/On** tags.

```
// @formatter:off
String createTable = "CREATE TABLE bicycles (\r\n" +
    "  manufacturer text,\r\n" +
    "  model text,\r\n" +
    "  front_cogs list<int>,\r\n" +
    "  rear_cogs list<int>,\r\n" +
    "  size int,\r\n" +
    "  year int,\r\n" +
    "  PRIMARY KEY (manufacturer, model)\r\n" +
    ")";
session.execute(createTable);
// @formatter:on
```

20. Implement a method for creating indices:

```
protected void createIndex(String columnFamily, String columnName) {
    String createIndex = String.format("CREATE INDEX %s_%s ON %s (%s)",
        columnFamily, columnName, columnFamily, columnName);
    session.execute(createIndex);
}
```

Invoke the method:

```
createIndex("bicycles", "size");
createIndex("bicycles", "year");
```

21. The insertion uses the builder pattern (http://en.wikipedia.org/wiki/Builder_pattern):

```
Insert insert = QueryBuilder.insertInto(KEYSPACE_NAME, "bicycles");
insert.value("manufacturer", "KTM");
insert.value("model", "Strada 1000");
insert.value("size", 57);
insert.value("year", 2014);
insert.value("front_cogs", Arrays.asList(53, 39));
insert.value("rear_cogs", Arrays.asList(12, 13, 14, 15, 17, 19, 21, 23, 25, 28));
session.execute(insert);
```

22. Run some queries:

```
System.out.println("Query 1");
for (Row row : resultSet1) {
    System.out.println(row);
    String manufacturer = row.getString("manufacturer");
    String model = row.getString("model");
    List<Integer> frontCogs = row.getList("front_cogs", Integer.class);
    List<Integer> rearCogs = row.getList("rear_cogs", Integer.class);
    System.out.println(manufacturer + " " + model + ", front cogs: "
        + frontCogs + ", rear cogs: " + rearCogs);
}

System.out.println("Query 2");
String select2 = QueryBuilder.select("manufacturer").from("bicycles").
    where(QueryBuilder.eq("size", 57)).getQueryString();
ResultSet resultSet2 = session.execute(select2);
for (Row row : resultSet2) {
    System.out.println(row);
}
```

23. To improve the readability of the code, add a static import. **Tip:** use the Ctrl+Shift+M, while the cursor is in the eq part of `QueryBuilder.eq`.

```
import static com.datastax.driver.core.querybuilder.QueryBuilder.eq;

[...]

String select2 = QueryBuilder.select("manufacturer").from("bicycles").
    where(eq("size", 57)).getQueryString();
```

2.3 Integration

1. Install Maven: set the M2_HOME environment variable and add M2_HOME/bin to the PATH environment variable.
2. Check if Maven is working properly:

```
mvn --version
```

3. The standard Maven command for initiating the *build* process is:

```
mvn clean install
```

4. The build may throw the following error: Fatal error compiling: tools.jar not found:
To fix this, set the JAVA_HOME environment variable from the JRE's directory to the JDK's directory.
5. The CassandraClient class may cause an error related to the Java version, e.g. diamond operator is not supported in -source 1.5. Add the following lines to the POM file:

```
<build>
  <plugins>
    <plugin>
      <artifactId>maven-compiler-plugin</artifactId>
      <configuration>
        <source>1.7</source>
        <target>1.7</target>
      </configuration>
    </plugin>
  </plugins>
</build>
```

Source: <http://stackoverflow.com/questions/14043042/compiling-java-7-code-via-maven>

6. Once the build succeeds, try to run the application.

```
java -jar target\hu.bme.mit.sysint-0.0.1-SNAPSHOT.jar
```

7. It throws the following error:

```
no main manifest attribute, in target\hu.bme.mit.sysint-0.0.1-SNAPSHOT.jar
```

Define the main class in the POM file:

```

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-jar-plugin</artifactId>
  <configuration>
    <archive>
      <manifest>
        <mainClass>hu.bme.mit.sysint.cassandraclient.App</mainClass>
      </manifest>
      <manifestEntries>
        <mode>development</mode>
        <url>${pom.url}</url>
      </manifestEntries>
    </archive>
  </configuration>
</plugin>

```

Source: <http://docs.codehaus.org/pages/viewpage.action?pageId=72602>

8. It throws the following exception:

```

Exception in thread "main" java.lang.NoClassDefFoundError: com/datastax/driver/core/Statement
    at hu.bme.mit.sysint.cassandraclient.App.main(App.java:13)
Caused by: java.lang.ClassNotFoundException: com.datastax.driver.core.Statement

```

The reason for this is that the dependencies are not available for the application. To solve this, we have to

(a) Copy the dependencies:

```

<plugin>
  <artifactId>maven-dependency-plugin</artifactId>
  <executions>
    <execution>
      <phase>install</phase>
      <goals>
        <goal>copy-dependencies</goal>
      </goals>
      <configuration>
        <outputDirectory>${project.build.directory}/lib</outputDirectory>
      </configuration>
    </execution>
  </executions>
</plugin>

```

The execution tag shows an error:

maven-dependency-plugin (goals "copy-dependencies", "unpack") is not supported by m2e.

Use the provided **Permanently mark goal copy-dependencies in pom.xml as ignored in Eclipse build.** fix. This adds a long <pluginManagement> tag.

(b) Add them to the classpath:

```

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-jar-plugin</artifactId>
  <configuration>
    <archive>
      <manifest>

```

```

        [...]
        <addClasspath>true</addClasspath>
        <classpathPrefix>lib/</classpathPrefix>
    </manifest>
</archive>
</configuration>
</plugin>

```

For some libraries (e.g. emfjson) you may get a `java.lang.NoClassDefFoundError` caused by a `java.lang.ClassNotFoundException`. The reason for this is that Maven sometimes resolves the SNAPSHOT string to the actual snapshot version in the generated JAR's MANIFEST.MF file (e.g. `org.eclipselabs.emfjson-0.7.0-SNAPSHOT.jar` becomes `org.eclipselabs.emfjson-0.7.0-20140221.135604-5`) but the filename of the dependency's JAR stays the same. To solve this, add a `<useUniqueVersions>>false</useUniqueVersions>` element between the `<manifest>` tags (source: <http://jira.codehaus.org/browse/MJAR-156>).

Source: <http://stackoverflow.com/questions/97640/force-maven2-to-copy-dependencies-into-target-lib>

9. The old slf4j warning is back. The solution is to create a the `src/main/resources` folder and move the `log4j.properties` file to it.
10. If you would like to package your project as one “fat JAR”, use the `maven-assembly-plugin`.

```

<plugin>
  <artifactId>maven-assembly-plugin</artifactId>
  <configuration>
    <archive>
      <manifest>
        <mainClass>hu.bme.mit.sysint.cassandraclient.App</mainClass>
      </manifest>
    </archive>
    <descriptorRefs>
      <descriptorRef>jar-with-dependencies</descriptorRef>
    </descriptorRefs>
  </configuration>
  <executions>
    <execution>
      <id>make-assembly</id> <!-- this is used for inheritance merges -->
      <phase>package</phase> <!-- bind to the packaging phase -->
      <goals>
        <goal>single</goal>
      </goals>
    </execution>
  </executions>
</plugin>

```

Source: <http://stackoverflow.com/questions/574594/how-can-i-create-an-executable-jar-with-dependencies-using-maven>

If you just want to get the fat JAR, build the project with:

```
mvn clean compile assembly:single
```

2.4 Version control

1. Start by cloning the repository to your local computer. To do this, go to the **Git Repositories** view (from the **Window | Show View | Other...** window). Provide the clone URI to clone the Git repository.

2. Right click the project and choose **Team | Share project...**. Choose **Git** and select the repository.
3. Add a `.gitignore` file to the project. Search for GitHub's `gitignore` repository (<https://github.com/github/gitignore>) repository and use the `Maven.gitignore` and the `Java.gitignore` files as a guideline.
4. **Commit** and **Push** the project.

2.5 Eclipse tips

- In the **Window | Preferences** dialog, go to **Java | Editor | Typing**.
 - In the **Automatically insert at correct position** group, check the **Semicolons** checkbox.
 - In the **In string literals** group, check **Escape text when pasting into a string literal**.